

Artificial Intelligence Summer School

Multiagent Decision Making

Laurent.Vercouter@insa-rouen.fr

From Artificial Intelligence...

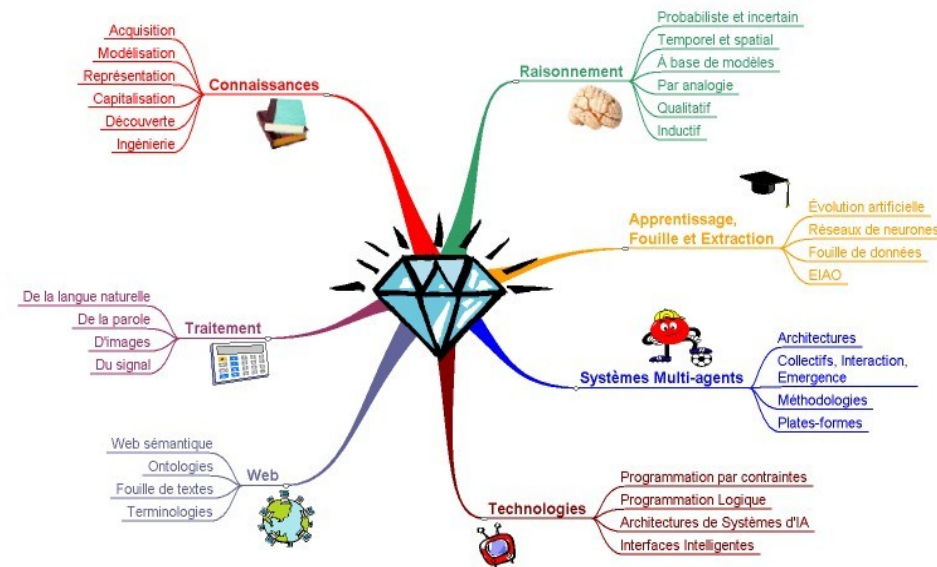
Creation of the field « Artificial Intelligence » at the Darmouth conference, 1956

2 approaches :

- Understanding and reproducing human intelligence
- Developing computer programs able to solve « complex problems »

« the construction of computer programs that perform tasks that are, for the moment, more satisfactorily accomplished by human beings, as they require high-level mental processes such as perceptual learning, memory organization and critical reasoning » (Marvin Minsky)

AI today (def. AFIA)



... to Distributed Artificial Intelligence...


In the 90's, the development of networks and multi-processor architectures paves the way for distributed and/or parallel systems.

Distributed Artificial Intelligence encompasses :

- Parallel A.I.
 - Improved computing performance
 - Languages for managing concurrency and parallelism
- Distributed problem solving
 - Problem decomposition for multiple solutions
 - Competing or complementary solutions
 - Need for synchronization of shared knowledge and partial solutions
- Multi-Agent Systems (MAS)

...to Multi-Agent Systems

A **MAS** is a set of agents that interact together and with a shared environment.



Macroscopic perspective: the MAS has a global behavior resulting from the actions and interactions of the agents that make it up.

Microscopic perspective: an agent is conceived as an intelligent and autonomous software entity.

- **Top-down** development: use expected macroscopic behavior as a starting point for developing agent models
- **Bottom-up** development: develop agent models and observe the resulting macroscopic effects

Agenda

- **Introduction to Multiagent systems**
 - Agent architectures
 - Main application domains
- Decision making
- Reactive agents
- Cognitive agents

Intelligent Agent (1)

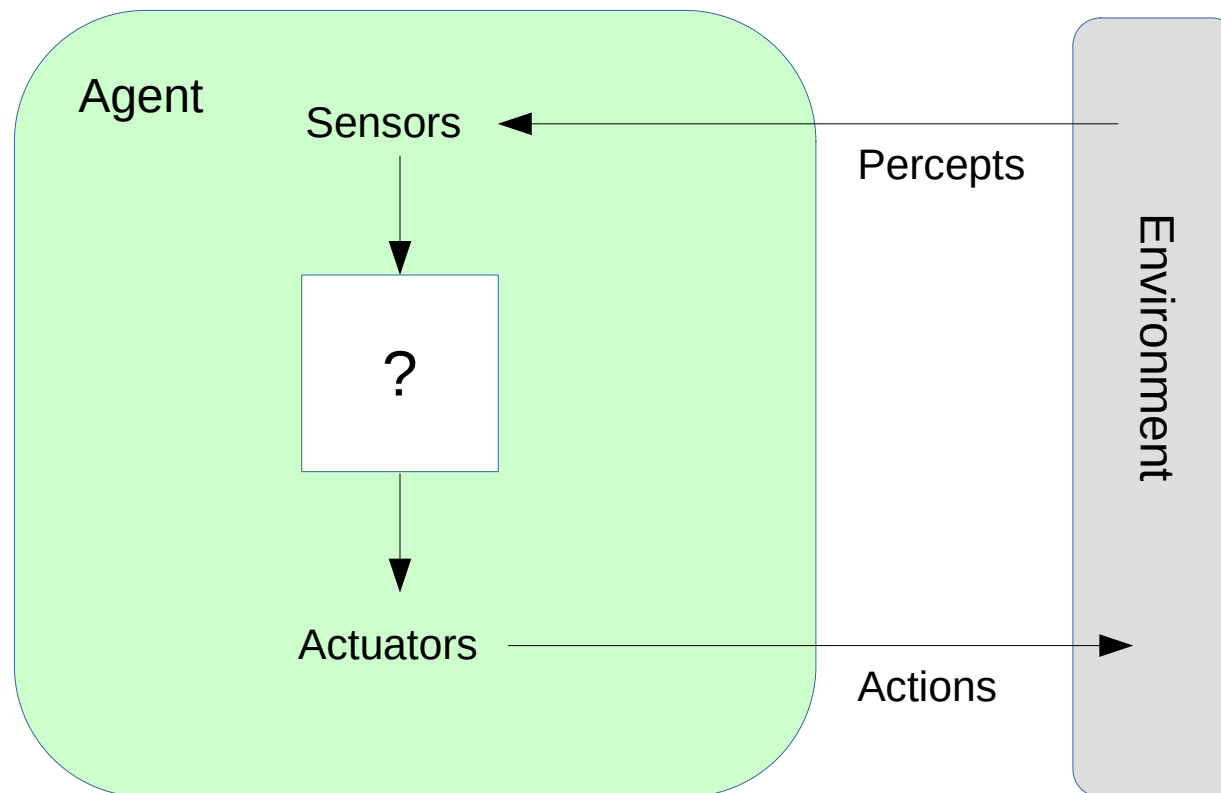
An agent is a physical or virtual entity (def. Ferber, 91) that

- is capable of acting in an environment,
- can communicate directly with other agents,
- is driven by a set of tendencies (in the form of individual objectives, etc.),
- has its own resources,
- is capable of perceiving its environment (albeit to a limited extent),
- possesses skills and offers services,
- ...

Intelligent Agent (2)

General scheme

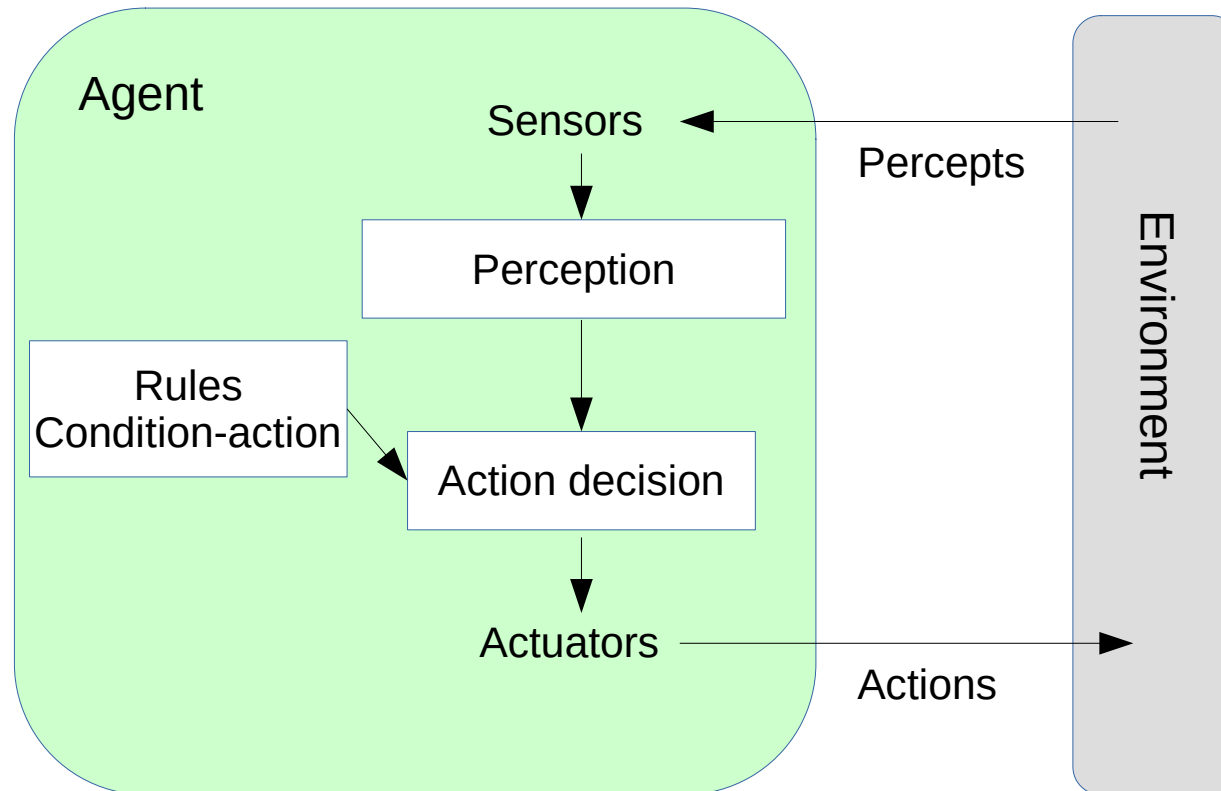
(source « *Artificial Intelligence* », Russel & Norvig, 3rd édition)



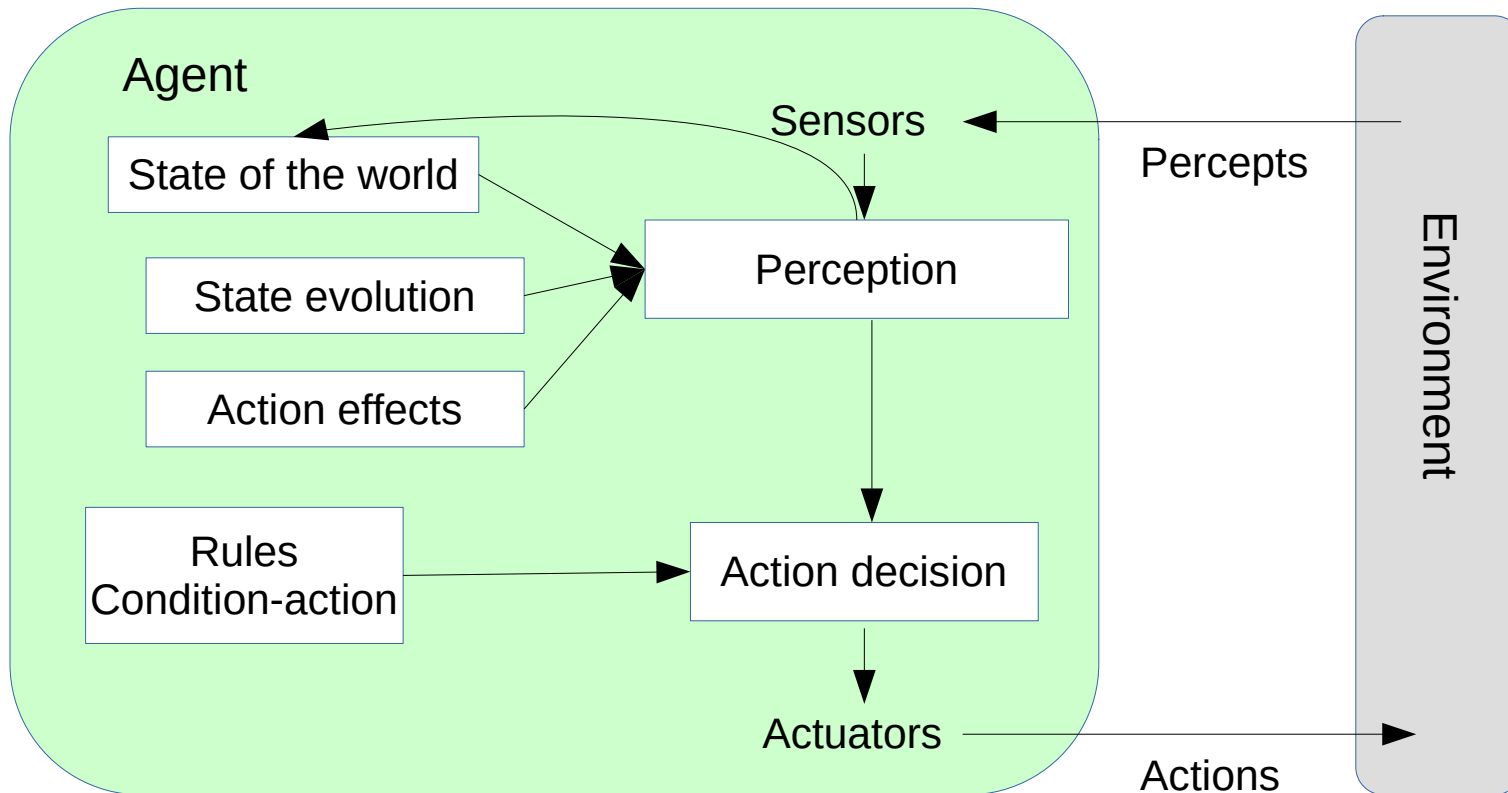
Simple reactive architecture

Example : thermostat reactive agent

- Perception: ***temperature*** = X
- Possible actions: ***heat***, ***turn off***
- Rules: if ***temperature*** < 19 then ***heat***, if ***temperature*** \geq 19 then ***turn off***



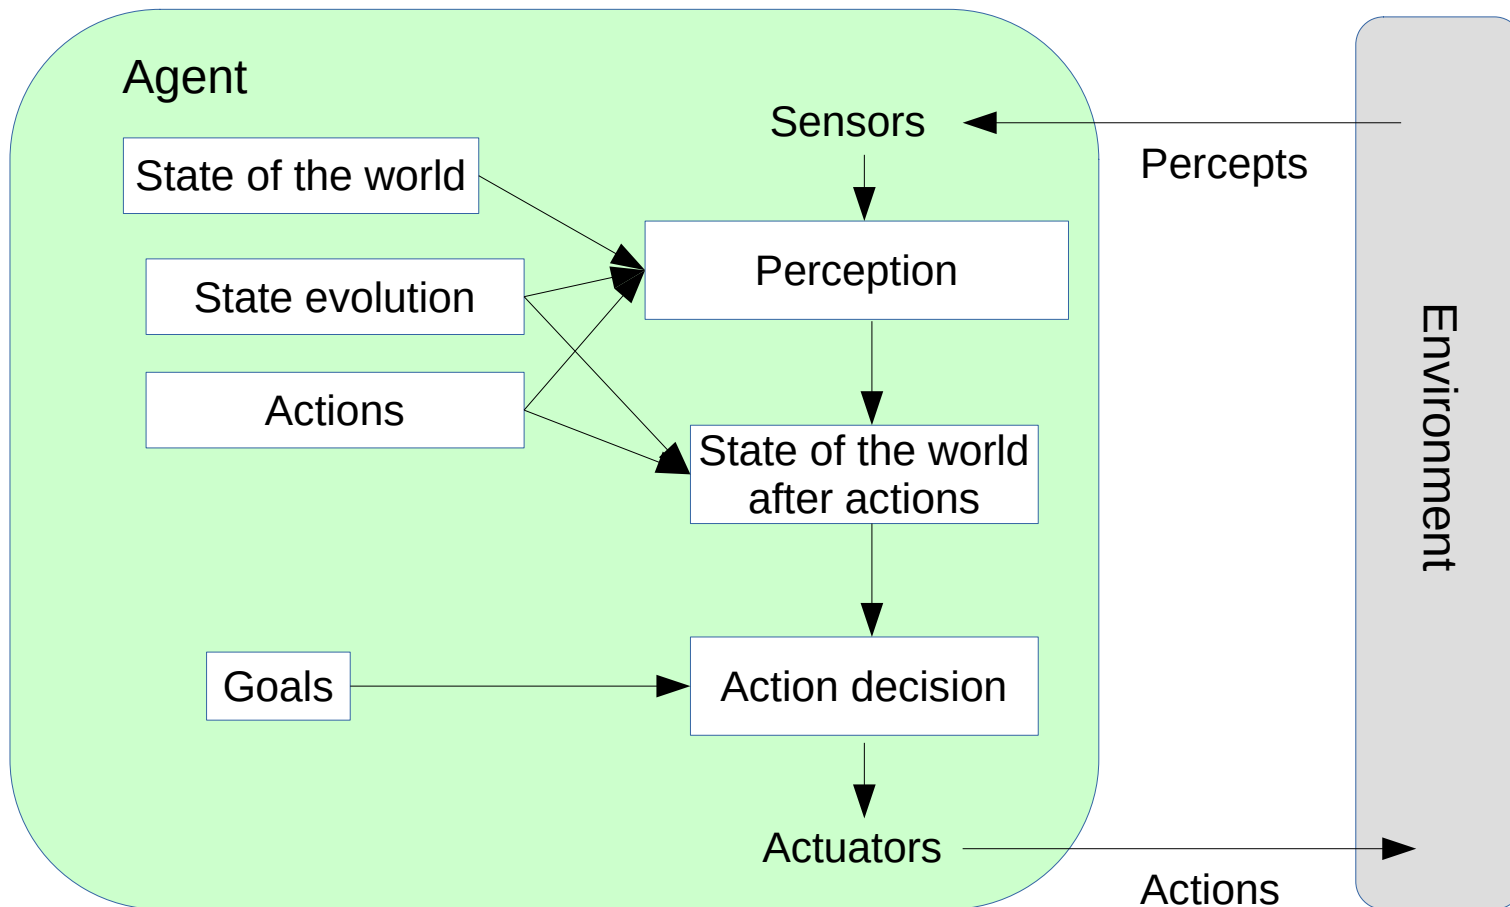
Reactive architecture with internal state



Cognitive architecture based on goals

Example : thermostat cognitive agent

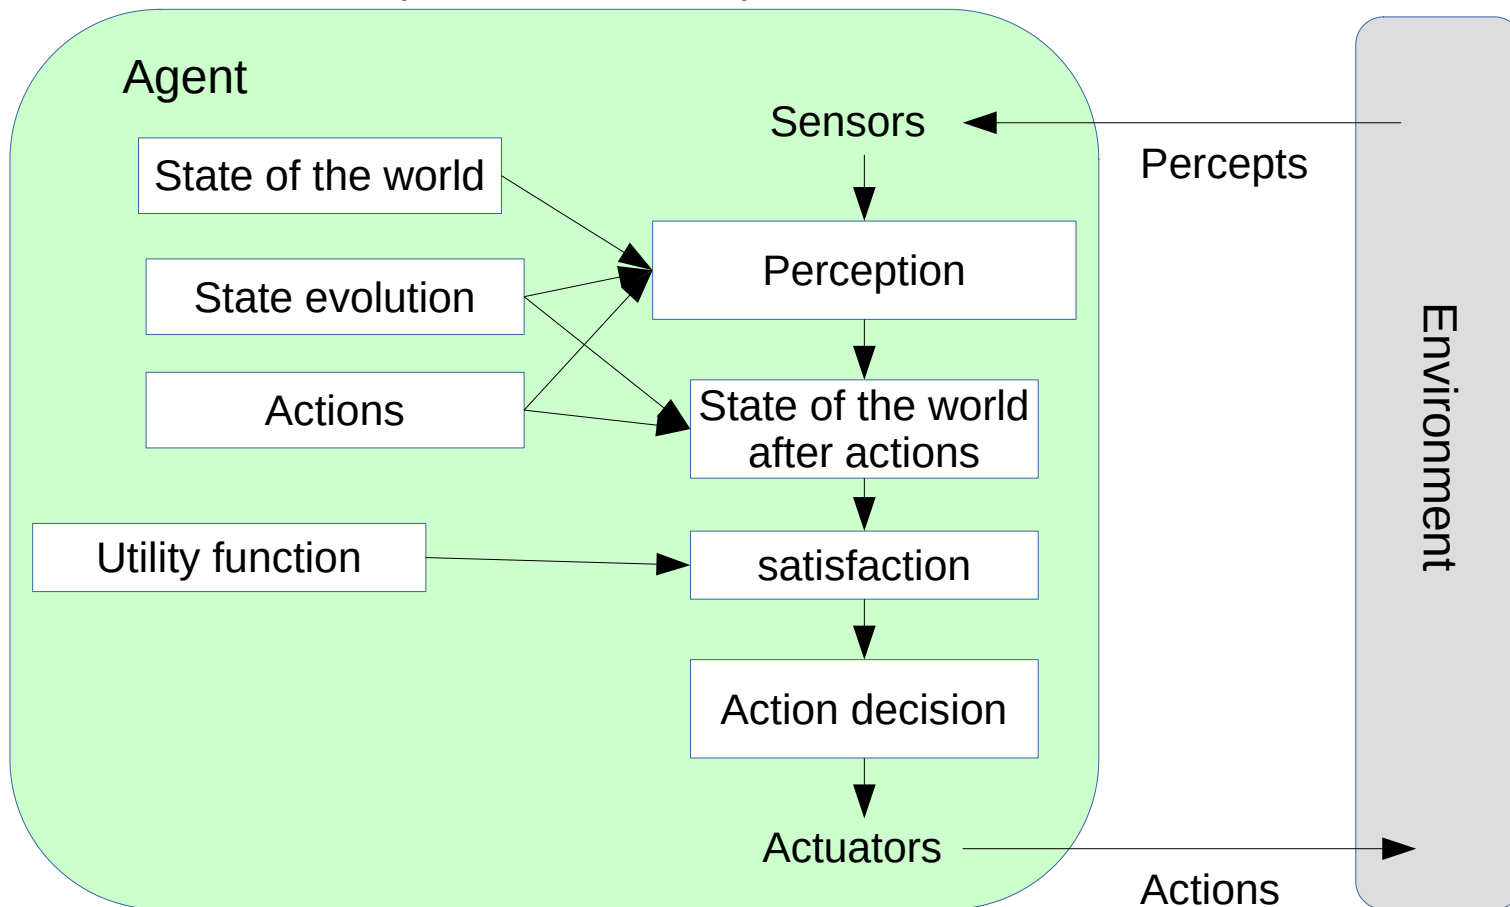
- State of the world: **temperature** = X , **temperature increasing**
- Possible actions: **heat** (effect = temperature++), **turn off** (effect = \emptyset)
- Goals: **temperature** ≥ 19



Cognitive architecture based on utility

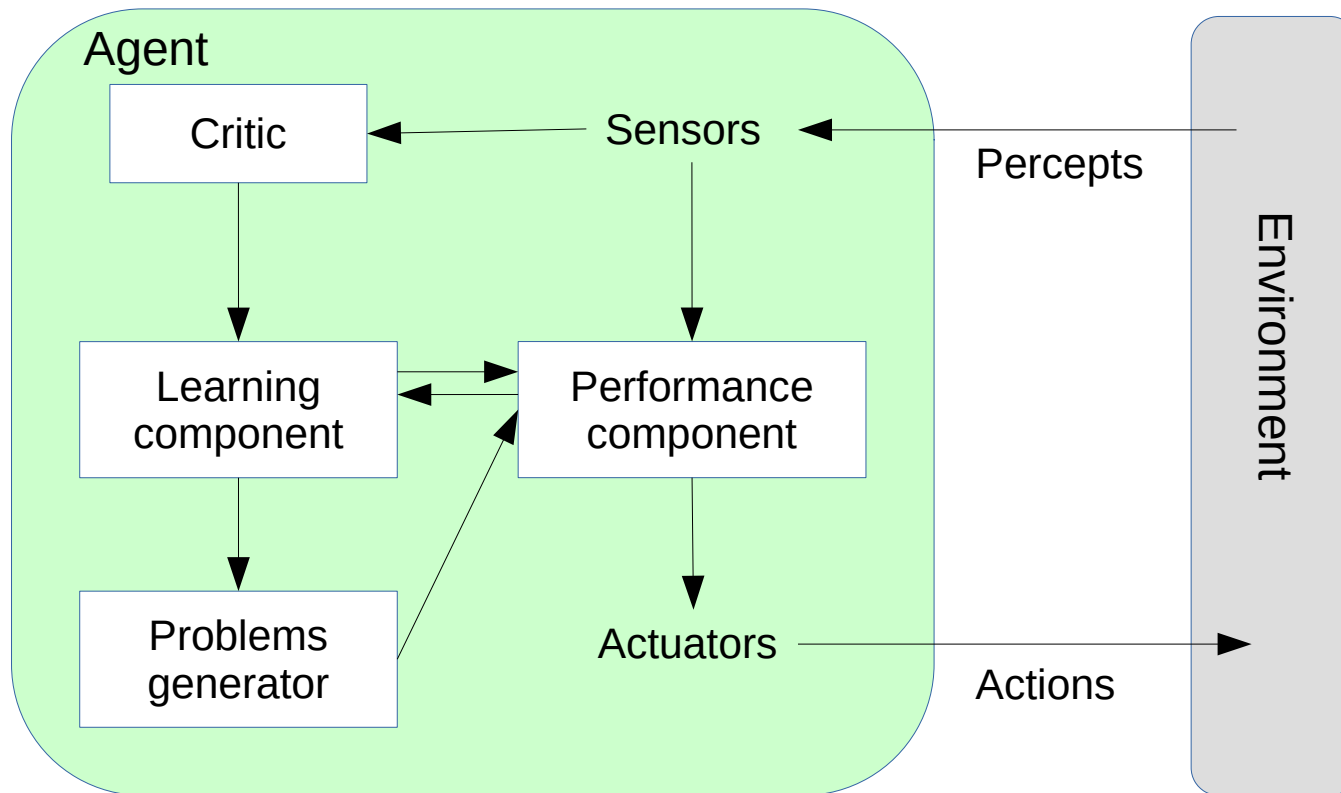
Example : thermostat cognitive agent

- State of the world: **temperature** = X
- Possible actions: **heat** (effect = $temperature++$), **turn off** (effect = \emptyset)
- Utility function: $minimize(|temperature-21|)$



Learning agent

The performance component groups here the components architecture seen previously.



Examples of application of an autonomous agent

- Autonomous robots



- Software personal assistant



- Conversational agent

- <https://www.youtube.com/watch?v=K6kcv3zwoo8>



Agent vs. Object

In many MAS programming languages, an agent is implemented as an object, but we are looking for the following properties:

- **Autonomy**
 - Internal: the agent's control mechanism is its own
 - External: no control over the behavior of other agents
- **Pro-activity**
 - Behavior directed by its own objectives (rules, goals, utility function, learning function)
- **Sociability**
 - Ability to interact with other agents

From one agent to a MAS

The co-existence of several agents in the same environment introduces a "social" dimension into the system.

- Dependencies between agents
 - Ex: A's actions can influence B's goals
- Interactions between agents
 - Ex: A communicates its goal to B to influence its behavior
- Social knowledge
 - Ex: A incorporates its perception of B's goals and knowledge into its own knowledge

Properties of a MAS

At the global level, we look for the following properties in the design of a MAS:

- System openness
 - During execution, agents can enter or leave the system
- Decentralized control
 - The control loop is implemented locally in each agent
- Agent heterogeneity
 - Agents with different motivations and functions can coexist


Application domains of MAS

- Distributed problem solving
 - Distributing AI algorithms
- Modeling and simulating autonomous entities
 - Reproduction or validation of individual and collective behavior models
- Distributed software applications
 - Service deployment in heterogeneous networks
- Interaction and communication with humans
 - Task delegation or assistance


Distributed problem solving

- Examples of spatial problem solving

Path finding



Cartography



<https://www.youtube.com/watch?v=8F6caMF7MXg>

- Characteristics
 - Actions to be chosen according to a local representation of an agent's world
 - Sequence of coordinated interactions between agents
 - Actions constrained by a simulated environment and agent position

Distributed software application

- Examples from collective robotics

Robocup



Robocup rescue



- Characteristics
 - Actions to be chosen according to a local representation of an agent's world
 - Communications between agents may be limited in scope
 - Highly uncertain perceptions and effects of actions

Agent-based simulation

- Example from situated individual-centric simulation

Traffic simulation



MASSIVE




- Characteristics
 - Actions to be chosen according to a local representation of an agent's world
 - Sequence of coordinated interactions between agents
 - Actions constrained by a simulated environment and agent position

Emergence and collective intelligence

Emergence

- In a hierarchical system of increasing complexity, a phenomenon or entity that originates at an earlier level is called emergent



Collective intelligence

- Inspiration from natural social systems (insect colonies, flocking, etc.)
- No overall coordination or supervision
- The fundamental principle of C.I. is that entities can produce functional behavioral patterns through direct or indirect interaction.

Agenda

- Introduction to Multiagent systems
- **Decision making**
 - Rationality
 - Basics of logic programming
- Reactive agents
- Cognitive agents

Rationality


- Rational thinking according to Aristotle (syllogism)
- In the 19th century, logicians defined formal statements at the intersection of philosophy and mathematics
- Rational action
 - Automatically defining the right action to perform
 - "right action" → existence of a goal or function to be optimized
- Limited rationality (H. Simon) = taking into account limited capacities or information

Automatic rational decision making

For software agents the « good » action to perform is defined by

- associating actions and situations
 - Reactive agents
- Defining goals and actions effects
 - Logical cognitive agents
- Defining several possible effects of agents
 - Probabilistic cognitive agents

A language for developing rational agents : Jason



- Interpreter for an extended version of the AgentSpeak language
- Paradigm of logic programming
- One of the most used platform to develop cognitive agents
- Downloadable in open source : <http://jason.sourceforge.net/>

Basics of first-order logic

In first-order logic, terms are made up of **variables** (whose symbol begins with a capital letter) and **functions** (in lower case).

- A variable is a term
- A function of arity n is a term, $f(t_1, \dots, t_n)$ with t_1, \dots, t_n being terms
- An arity 0 term is a constant, e.g. f

Example (to be used when simulating a vacuum cleaner robot)

- $pos(X)$ a pos function designates the robot's position X
- $dirty$ a constant indicating that the floor is dirty

Horn clauses

An Horn clause is written

$$A \text{ :- } B_1, \dots, B_n$$

with A, B_1, \dots, B_n being terms

An Horn clause specifies that if B_1, \dots, B_n are true then A is true.

If $n=0$ the Horn clause is a fact (always true), otherwise it is a rule.

Example

- *dirty* :- *pos*(*X*), *dirty*(*X*)

Logic programming

A logic program is a sequence of Horn clauses

For logical deduction or constraint satisfaction

- Expression of domain semantics by rules and facts
- Expression of an initial knowledge base by facts
- Deduction of new facts by logical inference

For rational decision-making

- Same possibilities for inference from facts
- Expression in the form of rules for actions to be taken when certain facts are true
- Updating a knowledge base by perceiving the effect of actions

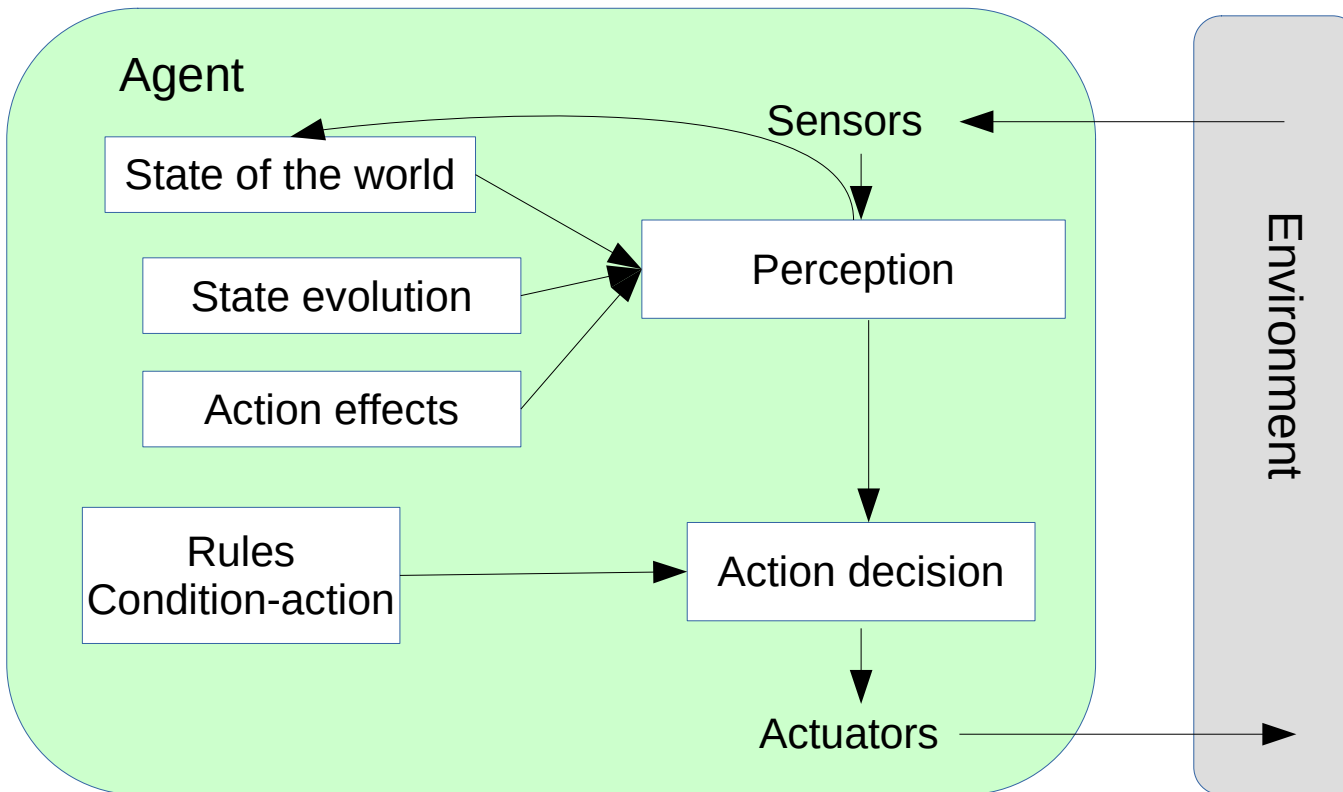
Example (written in Jason)

+dirty <- *suck* if fact *dirty* is true, action *suck* is rational

Agenda

- Introduction to Multiagent systems
- Decision making
- **Reactive agents**
 - General architecture
 - Programming reactive agents in Jason
- Cognitive agents

Jason reactive architecture



- Closed control loop: stimulus → response
- The environnement directly leads agents' behaviors
- No explicit world representaion
- No history

Reactive control function

```
do {  
    percepts:= see() ;  
    state:= interpret_inputs(percepts) ;  
    rule:= match(state,rules) ;  
    execute(rule[action]) ;  
} while (1) ;
```

Jason : Beliefs

Beliefs are represented by first-order logic literals

- *literal(term₁, ..., term_n) [annot₁, ..., annot_m]*

Examples

- *red(box1)[source(percept)].*
- *friend(bob,alice)[source(bob)].*
- *~dirty[source(self)]*

Jason : Beliefs dynamics

The **source** annotation is created automatically according to the origin of the belief

- Intention to add a belief : the + and – operators add/remove a belief coming from the agent itself
 - `+dirty ; // adds dirty[source(self)]`
 - `-dirty ; // removes dirty[source(self)]`
- By communication : effect of the send action
 - `.send(tom,tell,dirty) ; // sent by bob`
 - `// adds dirty[source(bob)] to tom's beliefs`
 - `.send(tom,untell,dirty) ; // sent by bob`
 - `// removes dirty[source(bob)] from tom's beliefs`

Jason : Description of a plan

An AgentSpeak plan has the following structure:

triggering_event : **context** \leftarrow **body**.

- **triggering_event** refers to the event that the plan has to consider
- **context** defines activation conditions
- **body** is the sequence of actions to perform

Jason : Triggering events

Triggering event can be from the following cases

- **+b** (belief addition)
- **-b** (belief removal)
- **+!g** (goal addition)
- **-!g** (goal removal)
- **+?g** (goal addition)
- **-?g** (goal removal)

Jason : Operators for context

Boolean operators

- `&` (and)
- `|` (or)
- `not` (not)
- `=` (unification)
- `>`, `>=` (relational)
- `<`, `<=` (relational)
- `==` (equals)
- `\ ==` (different)

Arithmetic operators

- `+` (sum)
- `-` (subtraction)
- `*` (multiply)
- `/` (divide)
- `div` (divide – integer)
- `mod` (remainder)
- `**` (power)

Jason : Negations

- Weak negation (not) = lack of belief
- Strong negation (\sim) = belief that a fact is false

+pos(X) : \sim dirty

<- move.

+pos(X) : not dirty & not \sim dirty

<- see(X).

Exercise

Write Jason plans to control a vacuum cleaner robot in a simulated environment with 4 squares.

Beliefs are automatically created by percepts

- *dirty* the robot square is dirty
- *clean* the robot square is clean
- *pos(X)* with $X = 1, 2, 3$ or 4 the robot is on the square X

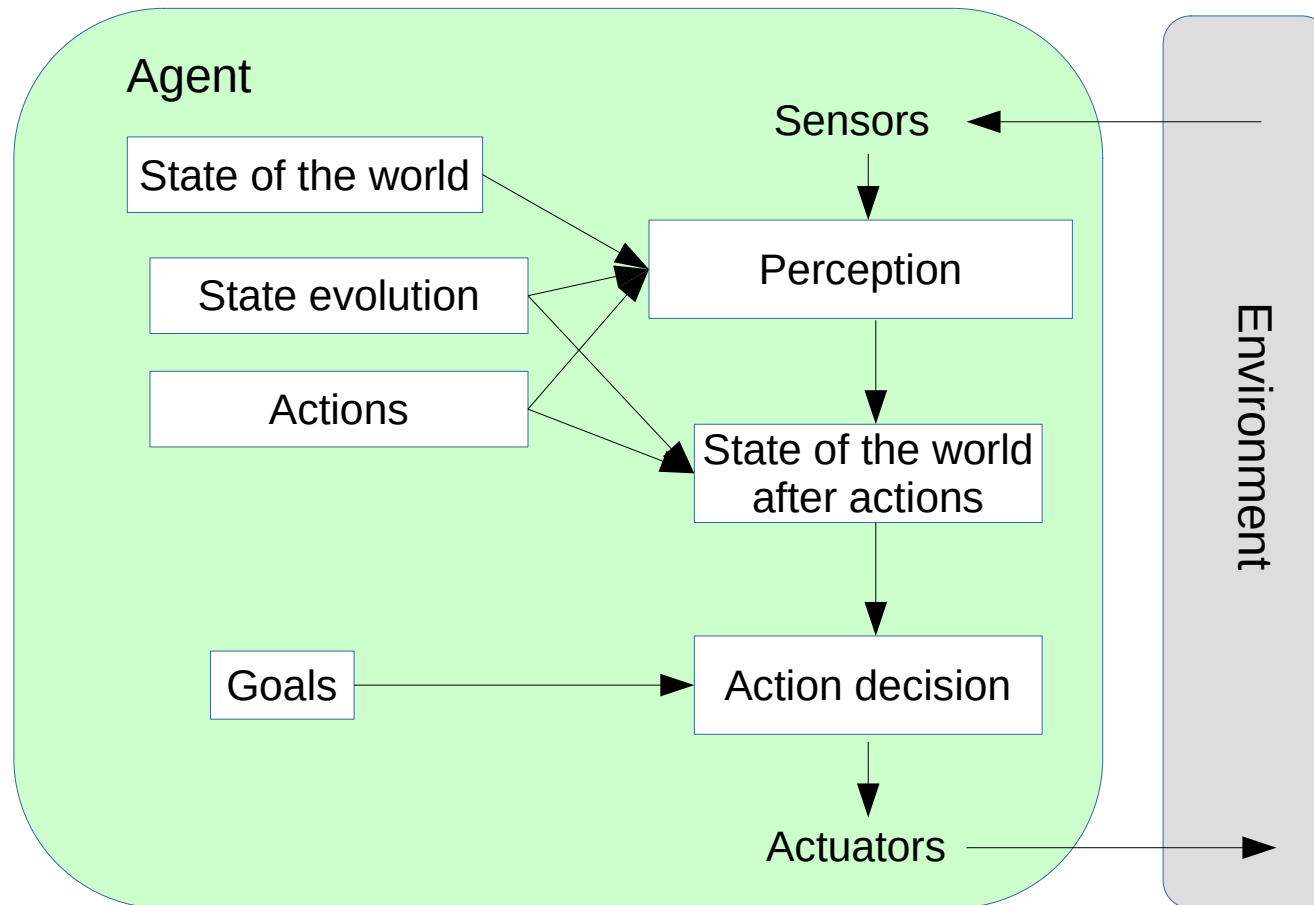
Possible actions

- *suck* the robot cleans the square
- *left* the robot moves left
- *right* the robot moves right
- *up* the robot moves up
- *down* the robot moves down

Agenda

- Introduction to Multiagent systems
- Decision making
- Reactive agents
- **Cognitive agents**
 - General architecture
 - Programming cognitive agents in Jason

Cognitive architecture



- Deliberative process between perception and action stages
- Explicit representation of the world
- Representation of past beliefs (history)

Cognitive control function

s : state

eq : event queue

s := initialize() ;

do {

 options:= option_generator(eq,s);

 selected:= deliberate(options, s) ;

 s := update_state(selected, s) ;

 execute(s) ;

 eq:= see() ;

} while (1) ;

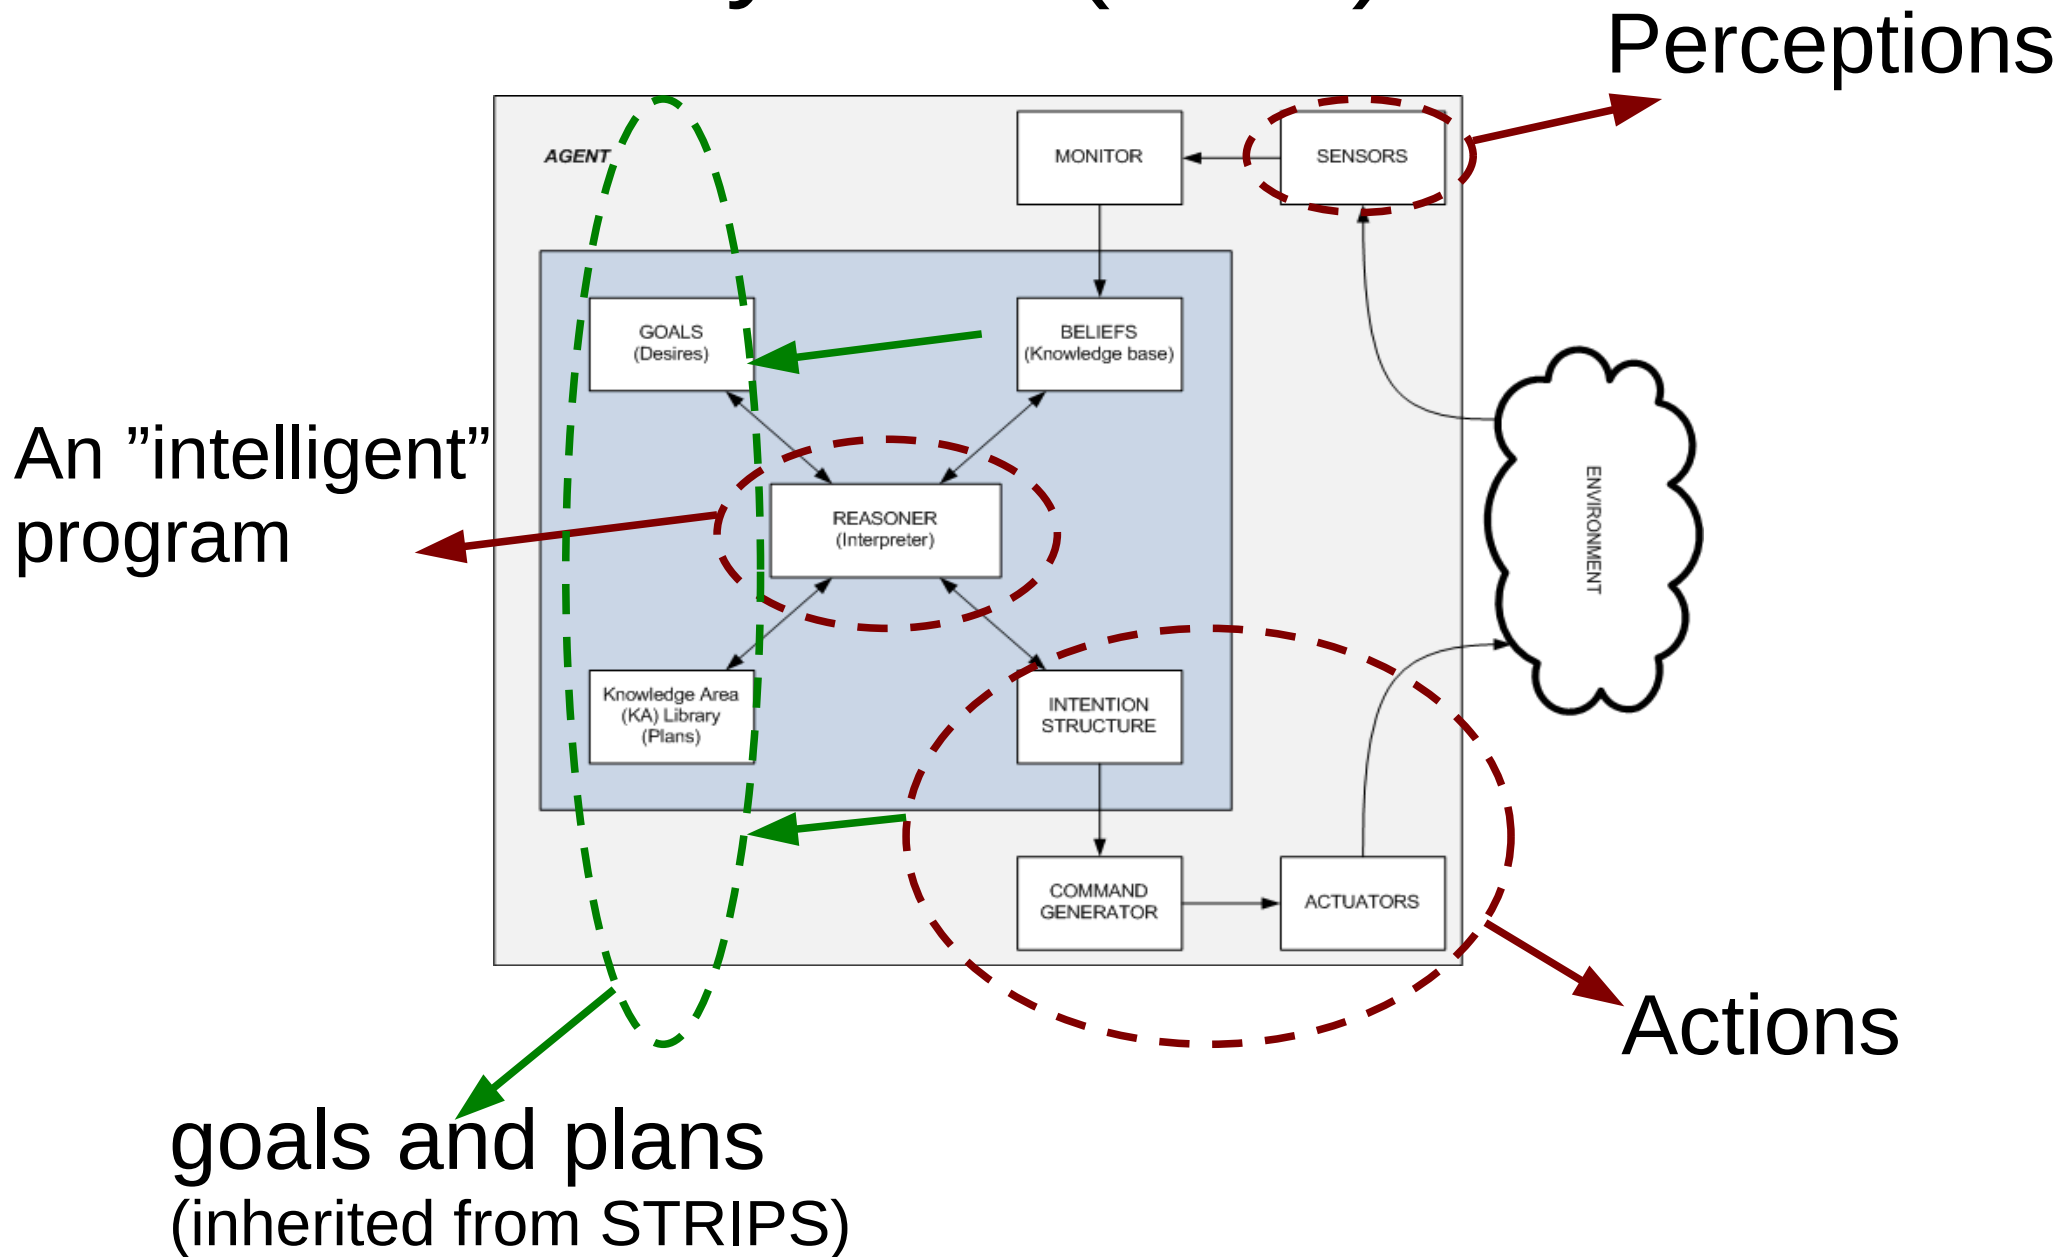
BDI model for cognitive agents

- BDI agent [Rao, Georgeff 91], according to this model an agent has :
 - **Beliefs** about itself and about the world,
 - **Desires** that may be contradictory,
 - A set of **Intentions** that are not in conflict,
 - A reasoning mechanism to update its beliefs, to choose a desire and to generate new intentions.

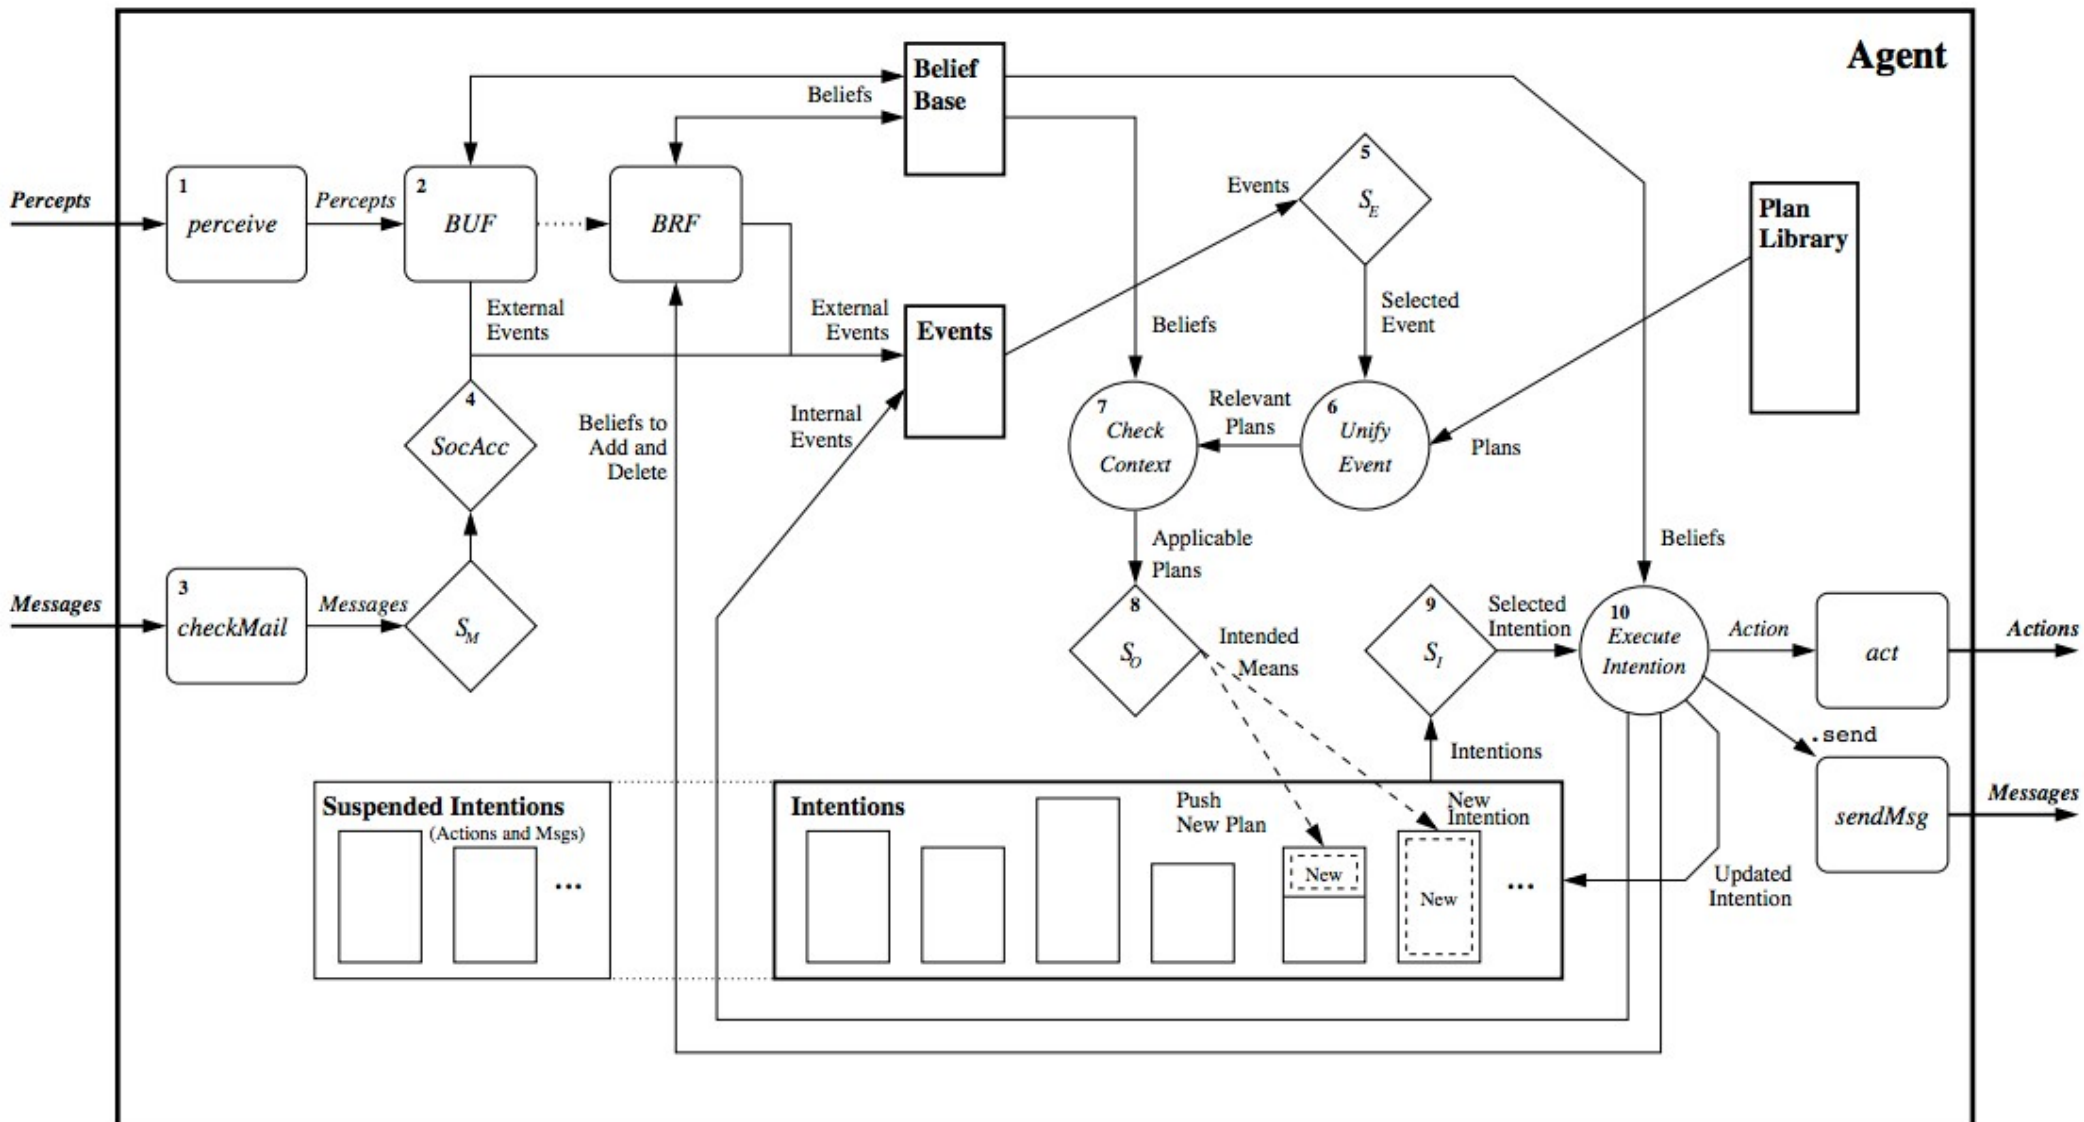
BDI languages and architectures

- BDI formal model is a reference for the development of several associated agents programming languages
 - AgentSpeak, dMars, Jack, Jadex, Jason, ...
 - Each of them complies (sometimes partially) to the BDI logics
an interpreter executes agents reasoning according to the BDI approach
- The architecture of a BDI agent defines
 - Its software components,
 - Their control processes,
 - Their data exchange processes,
 - The exchanges with entities external to the agent.

Example : Procedural Reasoning System (PRS)



The Jason reasoning cycle



Main elements of the language

Beliefs: represent the information available to an agent

Goals: represent the states the agent wants to achieve

Plans: represent the agent's "know-how", i.e. the means to achieve goals.

Events: represent changes in the agent's beliefs or goals

Intentions: instantiated plans for achieving goals

Main components of the architecture

Belief base: where beliefs are stored

Event set: where events that the agent must handle are stored

Plan library: where plans known to the agent are stored

Intention set: where the actions the agent must perform as part of a plan to achieve a goal are stored

Jason interpreter

- Perceives the environment and updates the belief base
- Processes new messages
- Selects events
- Selects **relevant** plans
- Selects **applicable** plans
- Creates/updates intentions
- Select intention to execute

Goals

- Achievement goal : a task that has to be accomplished
 - !goal
- Test goal : a goal to know an information
 - ?goal
- Examples
 - !write(book) ; // adds the goal !write(book)[source(self)]
 - ?publisher(P); // adds the goal ?publisher(P)[source(self)]

Goals dynamics (1)

The **source** annotation is automatically created according to the origin of the goal

- Intention to add a goal : the plan operators ! and ? adds/removes a goal coming from the agent itself
- By communication : effect of the send action for achievement goals
 - `.send(tom,achieve,write(book)) ; // sent by bob`
 - `// adds write(book)[source(bob)] to tom's goals`
 - `.send(tom,unachieve,write(book)) ; // sent by bob`
 - `// removes write(book)[source(bob)] from tom's goals`

Goals dynamics (2)

- By communication : effect of the send action for test goals
 - `.send(tom,askOne,published(P),answer) ; // sent by bob`
 - `// adds ?published(P)[source(bob)] to tom's goals`
 - `// the response will use answer`
 - `.send(tom,askAll,published(P),answer) ; // sent by bob`

Body of a plan

The body of a plan can contain

- Operators on beliefs + - -+
- Operators on goals ! ? !!
- Actions

```
+rain : time to leave(T) & clock.now(H) & H >= T  
  <- !g1; // new sub-goal  
    !!g2; // new goal  
    ?b(X); // new test goal  
    +b1(T-H); // add belief  
    -b2(T-H); // remove belief  
    -+b3(T*H); // update belief  
    close(door). // external action
```

Plans examples

```
+green_patch(Rock)[source(percept)]  
: not battery charge(low)  
<- ?location(Rock,Coordinates);  
    !at(Coordinates);  
    !examine(Rock).
```

```
+!at(Coords)  
: not at(Coords) & safe path(Coords)  
<- move towards(Coords);  
    !at(Coords).
```

```
+!at(Coords)  
: not at(Coords) & not safe path(Coords)  
<- ...
```

```
+!at(Coords) : at(Coords).
```

Plans dynamics

The plan library is composed by

- plans initially defined by the developer
- plans dynamically modified by the actions
 - .add_plan
 - .remove_plan
- plans communicated by messages
 - tellHow
 - untellHow

Jason library

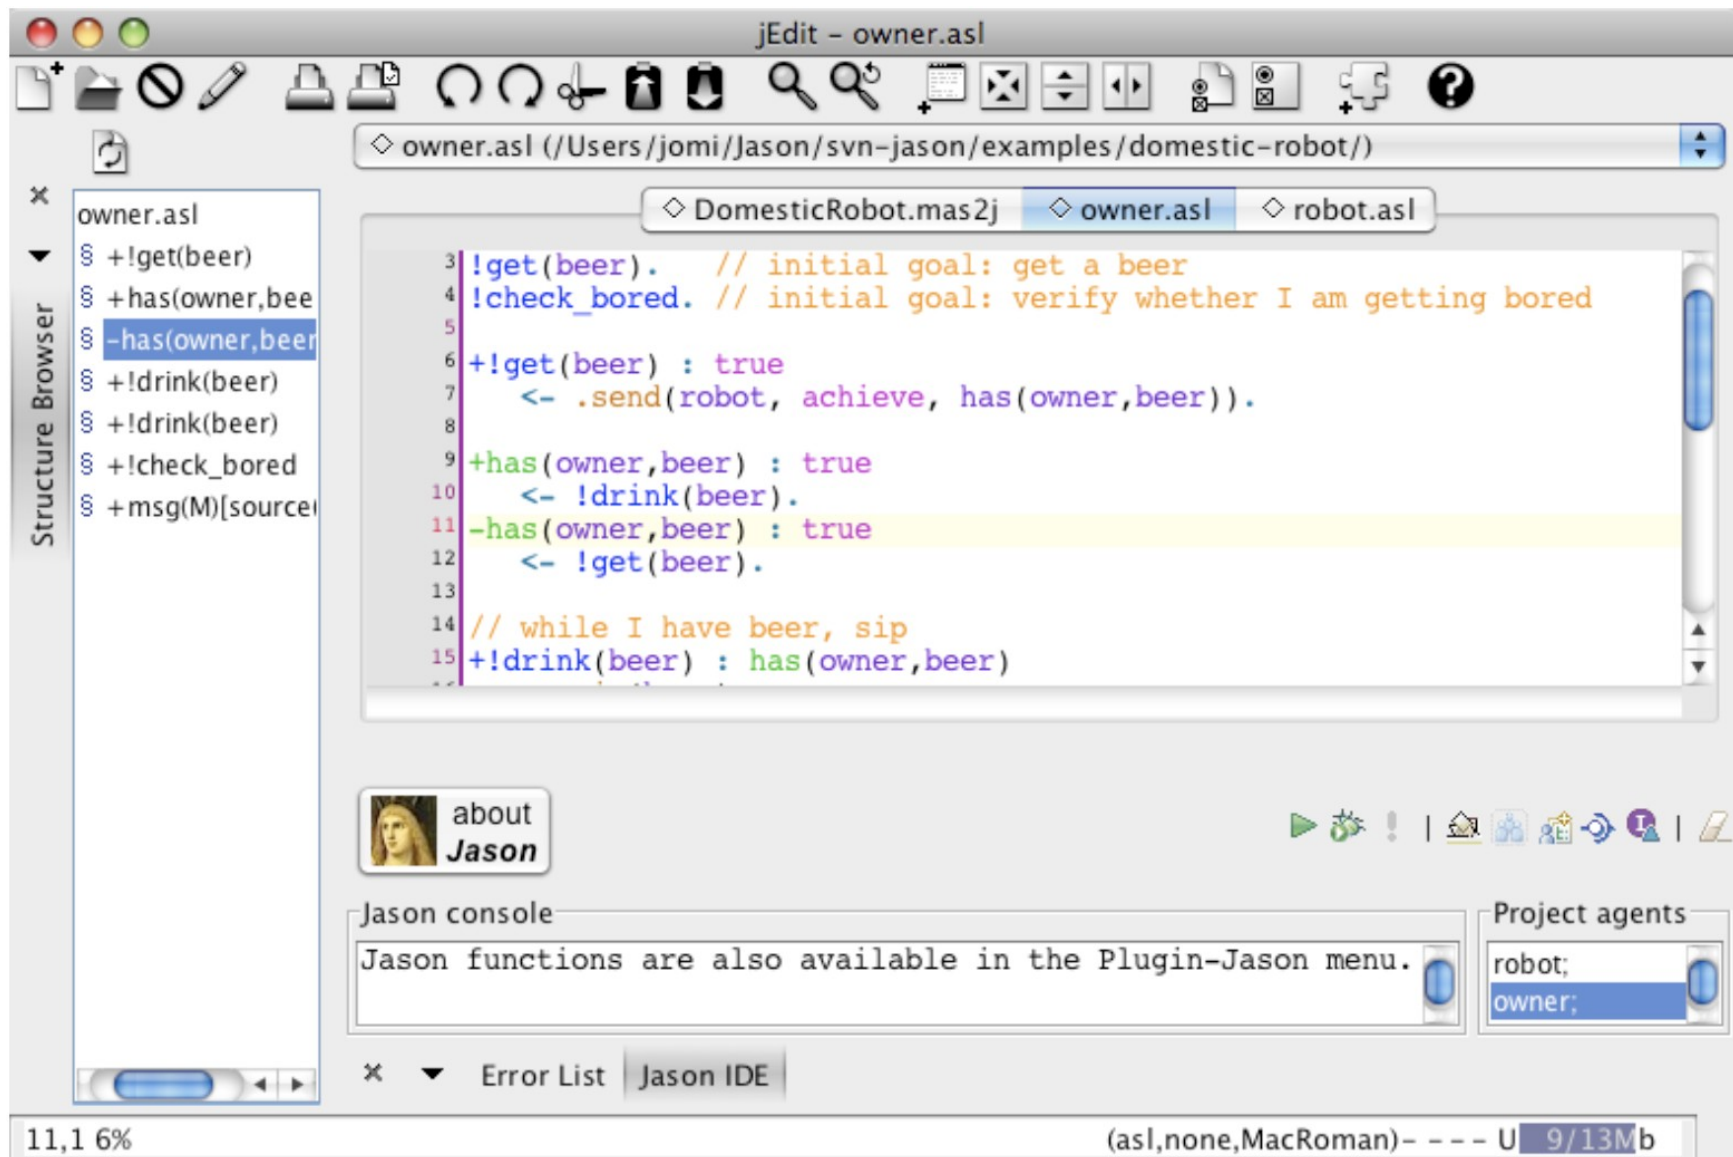
- Java library
 - Agent class extendable and able to interpret the Jason AgentSpeak language
 - Environment class to extend its own simulated environment
 - Multiagent communication infrastructure
- Definition of the MAS to execute in a configuration file

Jason configuration file

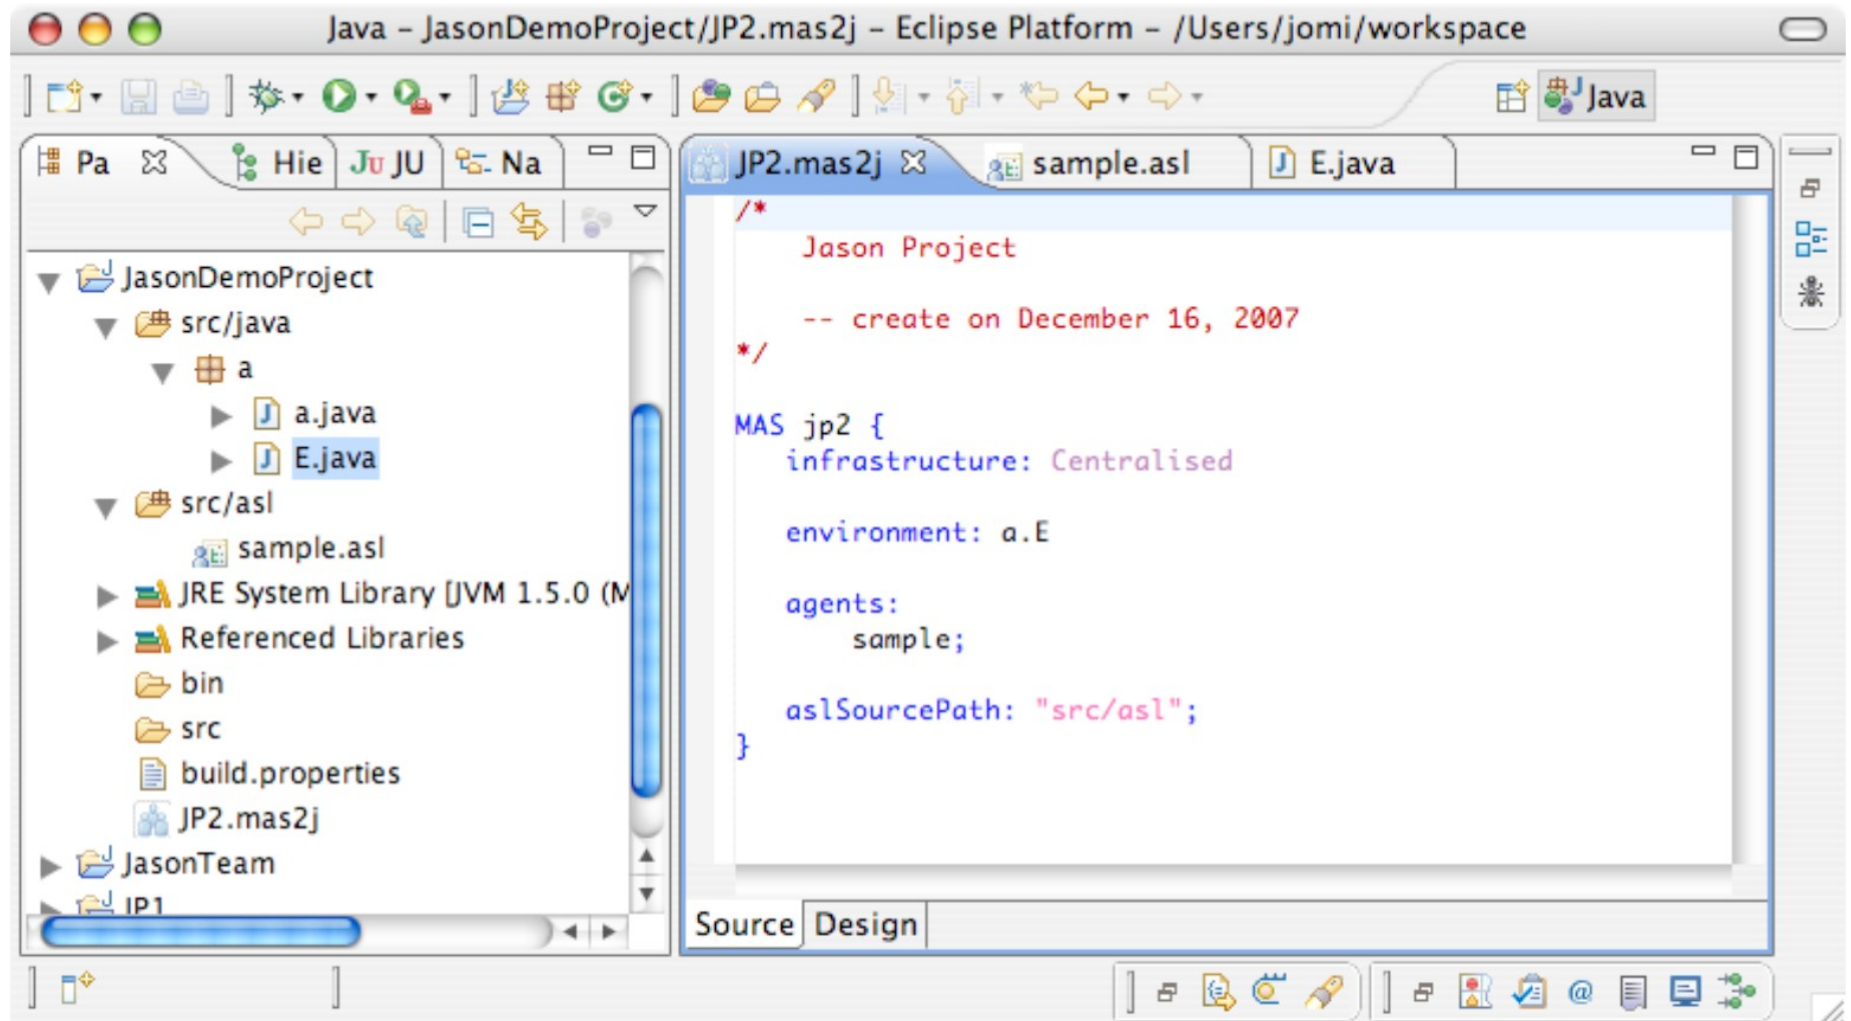
Example of configuration :

```
MAS my_system {  
  infrastructure: Jade  
  environment: robotEnv  
  agents:  
    c3po;  
    r2d2 at jason.sourceforge.net;  
    bob #10; // 10 instances of bob  
  classpath: "../lib/graph.jar";  
}
```


Plugin JEdit



Plugin Eclipse



Mind inspector

The screenshot displays the Jason Mind Inspector interface. The title bar reads ":: Jason Mind Inspector :: cycle 22 ::". On the left, the "Agents" panel lists "r2" and "r1", with "r1" selected. The main "Agent Inspection" panel shows the "Inspection of agent r1 (cycle #12)".

- Beliefs

- `pos(back,3,0)`_[source(self)]
- `pos(r1,3,0)`_[source(percept)]
- `pos(r2,3,3)`_[source(percept)]
- `garbage(r1)`_[source(percept)]

- Events

Sel	Trigger	Intention
X	<code>+!ensure_pick(garb)</code>	4

+ Options

- Intentions

Sel	Id	Pen	Intended Means	Stack (show details)
X	4		<code>+!ensure_pick(S)</code> <code>+!take(S,L)</code> <code>+!carry_to(R)</code> <code>+garbage(r1)</code> _[source(percept)]	<code>{ S = garb }</code> <code>{ S = garb, L = r2 }</code> <code>{ R = r2, Y = 0, X = 3 }</code>

Actions

Pend	Feed	Sel	Term	Result	Intention
X		X	<code>pick(garb)</code>	false	4

Agent History

A timeline slider at the bottom shows the simulation progress from Cycle 0 to Cycle 22, with a blue diamond marker positioned at Cycle 12.

At the bottom of the window, there is a "Run" button, a text field for "5" cycle(s) for, a dropdown menu set to "all agents", and a "view as:" dropdown menu set to "html".

Sources et liens

Livres de références

- Multiagent systems, G. Weiss, MIT Press, 2nd edition
- Multiagent Systems, M. Wooldridge, 2nd edition
- Les Systèmes Multi-Agents, J. Ferber
- Intelligence Artificielle, Russel & Norvig, 3ème édition

Cours accessibles en ligne

- O. Boissier (Ecole des Mines de Saint-Etienne)
<http://www.emse.fr/~boissier/enseignement/maop14/index.html>
- R. Courdier (université Réunion)
<http://lim.univ-reunion.fr/staff/courdier/cours/index.html>
- J.-P. Sansonnet (université Orsay) <http://perso.limsi.fr/jps/>
- ...